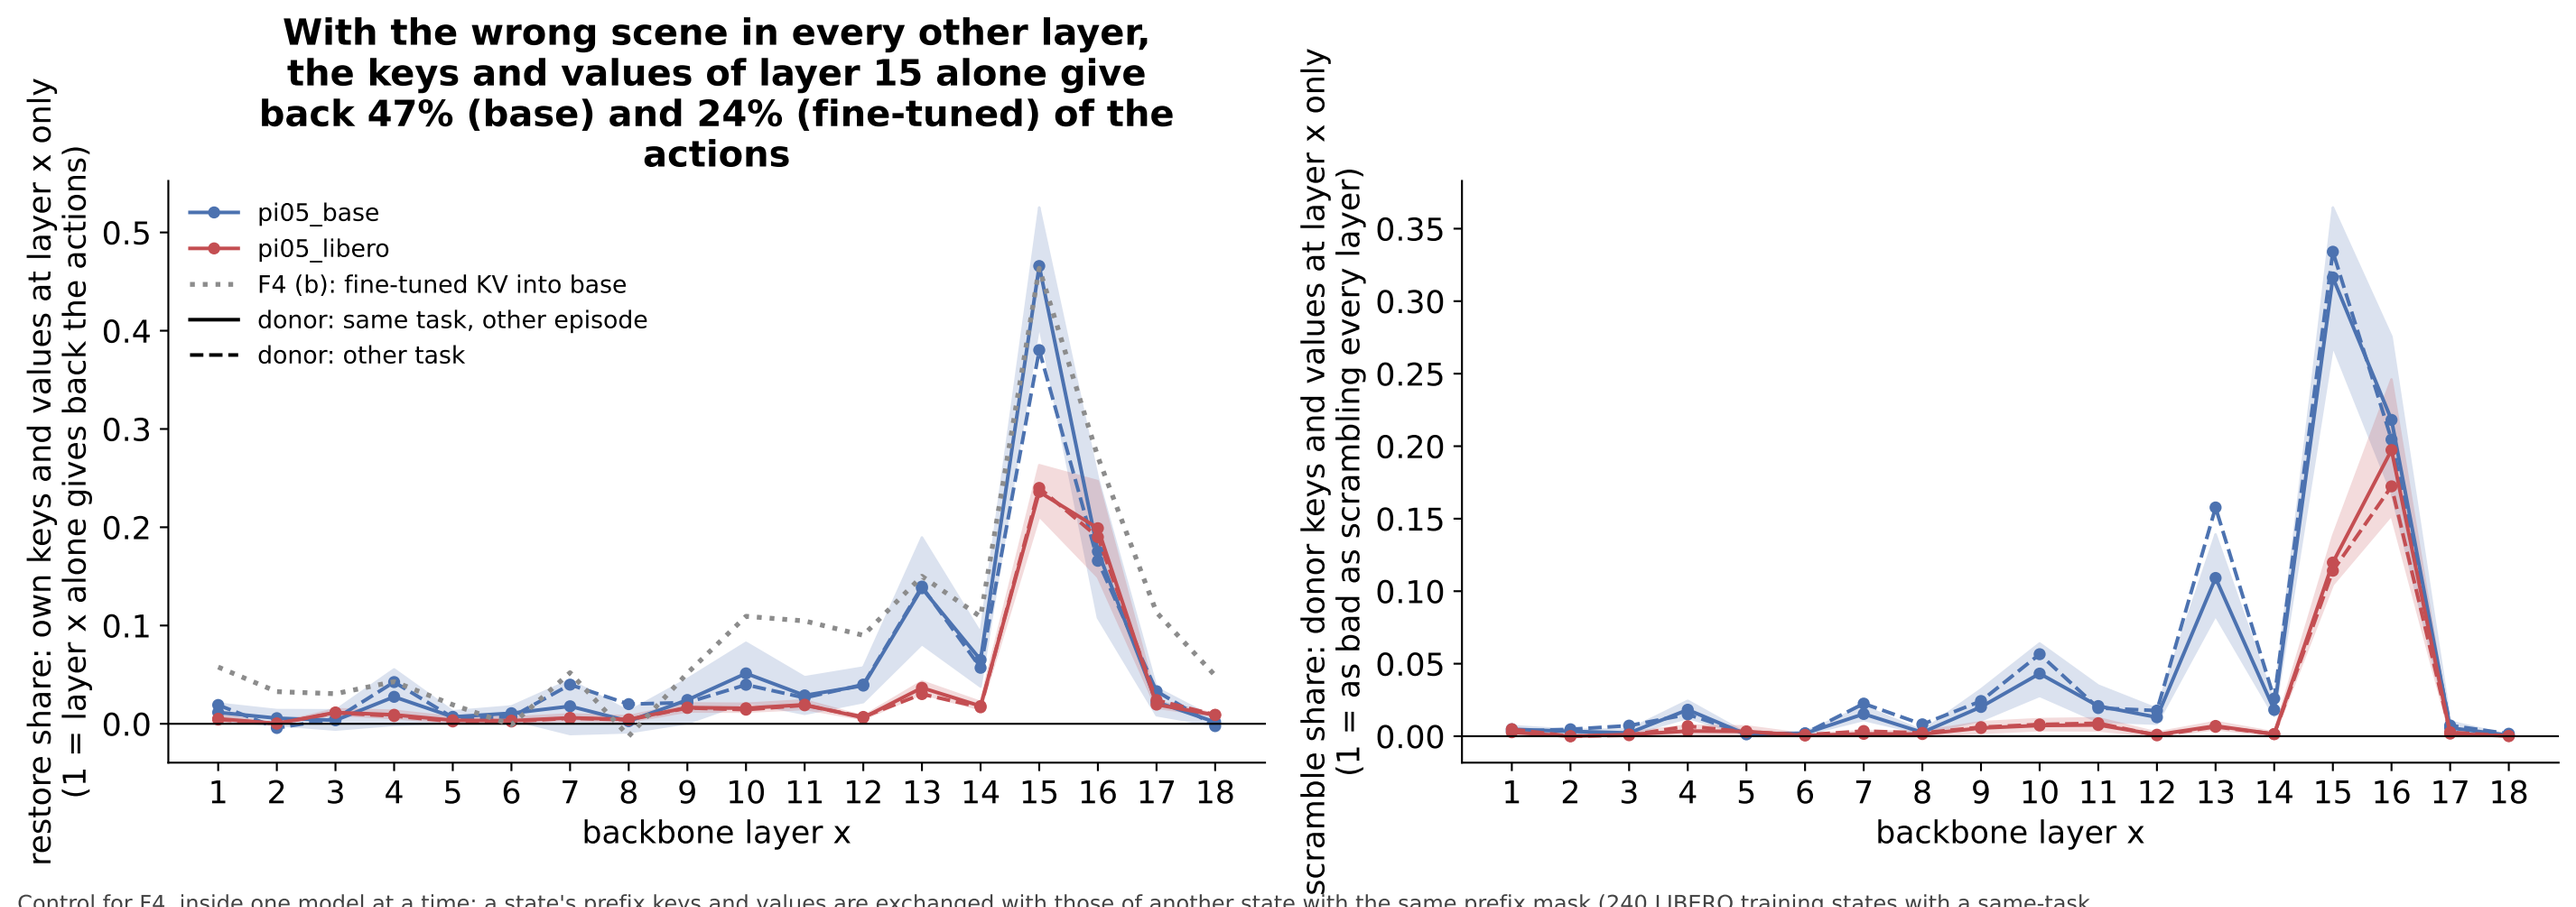

Control for F4, inside one model at a time: a state's prefix keys and values are exchanged with those of another state with the same prefix mask (240 LIBERO training states with a same-task donor, 228 with an other-task donor). Error is measured to the same model's own unpatched actions; bands are 95% bootstrap intervals over states (same-task donors). This shows which layers the action expert reads the scene from, not what fine-tuning changed. Pilot: pi05_base vs pi05_libero, which was fine-tuned on the full LIBERO data. This is a large fine-tune; the few-demo regime is left for later.